

CSE 156/L

Socket Overview

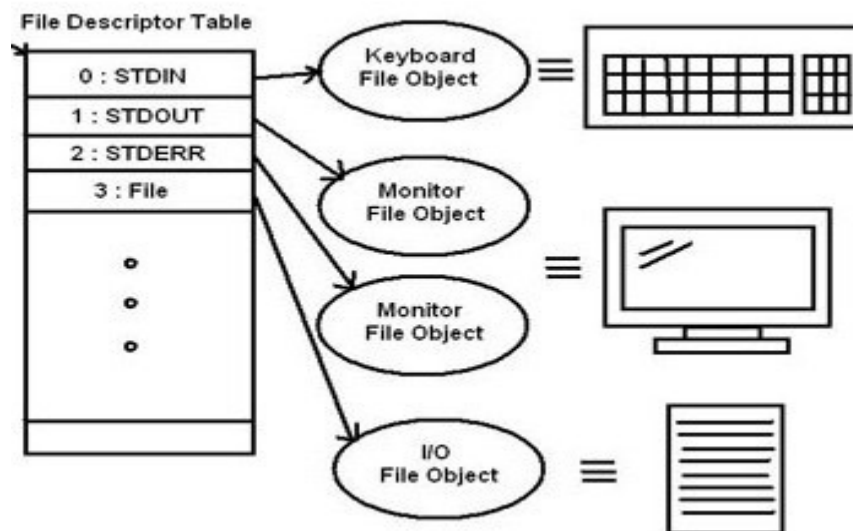
Motivation

- ❑ Developing network applications requires an Application Program Interface (API)
- ❑ TCP/IP protocol specifications do not include an API definition
- ❑ Many APIs defined by OS platforms
 - ❖ Sockets (BSD Unix), standardized by IEEE (POSIX)
 - ❖ TLI/XTI (AT&T System V)
 - ❖ Winsock (Microsoft)
 - ❖ MacTCP (Apple)

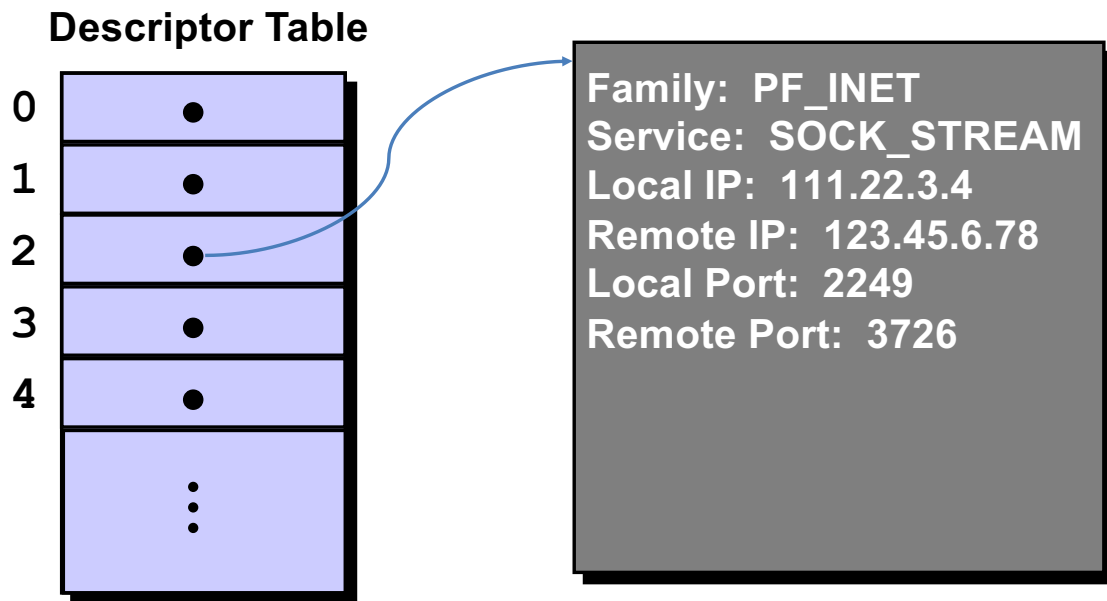
Typical Functions Needed

- ❑ Specify the endpoints of a network connection (local and remote)
- ❑ Open a connection (active open)
- ❑ Wait for an incoming connection (passive open or listen)
- ❑ Terminate a connection
- ❑ Abort a connection because of errors

Unix File Descriptor Table



Socket Descriptor Data Structure



POSIX Data Types

<code>int8_t</code>	signed 8-bit integer
<code>uint8_t</code>	unsigned 8-bit integer
<code>int16_t</code>	signed 16-bit integer
<code>uint16_t</code>	unsigned 16-bit integer
<code>int32_t</code>	signed 32-bit integer
<code>uint32_t</code>	unsigned 32-bit integer

Regular C data types: `u_char`, `u_short`, `u_int`, `u_long`

Creating a Socket

```
int socket(int family, int type, int proto);
```

- ❑ *family*: specifies the protocol family
 - ❖ AF_INET for TCP/IPv4, AF_INET6 for TCP/IPv6
 - ❖ PF_INET: practically socket.h defines PF_* equal to AF_*
- ❑ *type*: type of service
 - ❖ E.g., SOCK_STREAM or SOCK_DGRAM
- ❑ *proto*: protocol to use
 - ❖ Usually is 0 = default for given family/type
- ❑ Returns integer socket descriptor (-1 on error)
- ❑ Allocates socket data structure but does **not** define endpoint addresses

Address Structure

POSIX data types for endpoint address:

sa_family_t	address family
socklen_t	size of structure
in_addr_t	IPv4 address
in_port_t	transport port number

TCP/IPv4 Socket Address Structure

POSIX data types for endpoint address:

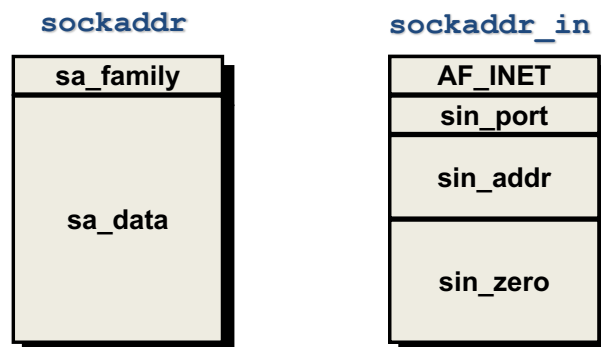
```
struct sockaddr_in {  
    uint8_t      sin_len;  
    sa_family_t  sin_family;    // uint8_t  
    in_port_t    sin_port;      // uint16_t  
    struct in_addr sin_addr;  
    char         sin_zero[8];  
};
```

```
// typedef uint32_t    in_addr_t;  
struct in_addr {  
    in_addr_t    s_addr;  
};
```

TCP/IP Socket Address Structure

- Generic socket structure (sockaddr) used in socket API prototypes (16 bytes)
- Cast protocol structures (e.g., sockaddr_in for IPv4) to this generic type

```
struct sockaddr {  
    unsigned short sa_family;    // address family, AF_XXX  
    char          sa_data[14];  // 14 bytes of protocol address  
};
```



Network Byte Order

- ❑ Network byte order for Internet protocols is big-endian
- ❑ IP address and TCP port number stored in `sockaddr` structure must be in network byte order
- ❑ Host byte order may be big- or little-endian
- ❑ Use library functions to convert between network and host byte order

Byte Order Conversion

‘h’ : host byte order
‘n’ : network byte order
‘s’ : short (16bit)
‘l’ : long (32bit)

```
uint16_t  htons(uint16_t);  
uint16_t  ntohs(uint16_t);  
uint32_t  htonl(uint32_t);  
uint32_t  ntohl(uint32_t);
```

Specifying Host Address

- ❑ The **bind()** system call is used to assign an address to an existing socket
 - I.e., IP address and port number of the local endpoint
- ❑ Typically used by servers to assign listening address and port.
- ❑ The servers have *well-known* port number assigned to the application
- ❑ Can also be used by clients to bind to a specific port
- ❑ *addrlen* is the sizeof *addr* specific structure

```
#include <sys/types.h>
#include <sys/socket.h>

int bind(int sockfd,
         const struct sockaddr *addr,
         socklen_t addrlen);
```

Specifying Host Address

- ❑ **bind()** returns 0 if successful or -1 on error
- ❑ Calling bind() with a port number of 0 results in the OS assigning an available port number
- ❑ How to determine IP address?
 - No general way to determine the IP address to bind to, when the host has multiple network interfaces
 - Common practice is to specify IP address as INADDR_ANY to let OS assign the IP address

```
#include <sys/types.h>
#include <sys/socket.h>

int bind(int sockfd,
         const struct sockaddr *addr,
         socklen_t addrlen);
```

bind() Example

```
int mysock, err;
struct sockaddr_in myaddr;

mysock = socket(AF_INET, SOCK_STREAM, 0);
myaddr.sin_family = AF_INET;
myaddr.sin_port = htons(portnum);
myaddr.sin_addr.s_addr = htonl(ipaddress);

err = bind(mysock,
           (struct sockaddr*) &myaddr,
           sizeof(myaddr));
```

TCP Client

Initiating a TCP Connection

```
int connect(int sockfd,
            const struct sockaddr *server,
            socklen_t addrlen);
```

- ❑ Clients initiate connection using **connect()** system call
 - ❖ Also sets up an endpoint address (IP address, port number) for the client socket
 - ❖ Clients don't need to call bind first
- ❑ Results in the OS performing TCP connection setup to remote host via 3-way handshake
 - ❖ *sockfd* is socket returned by **socket()** call
 - ❖ *server* contains the IP address and TCP port number of the server
 - ❖ Returns 0 normally, -1 on error

Closing a Connection

- ❑ Either side can initiate closing the connection with the **close()** system call
- ❑ Must be called by both sides to terminate cleanly
- ❑ If the other end has closed the connection, and there is no buffered data, reading from a TCP socket returns 0 to indicate EOF

Receiving Data from a Socket

```
int read(int fd, char *buf, int max);
```

max is the maximum number of bytes the read is willing to accept (size of buffer)

- By default, **read()** will block until data is available
 - Non-blocking option available
- Returns number of bytes read into the buffer
 - Between 1 and *max* bytes when connection is open
 - 0 indicates EOF (i.e., other side closed connection)

Sending Data to a Socket

```
int write(int fd, char *buf, int num);
```

num is the number of bytes to be written

- Writes can be made non-blocking
 - Might not be able to write all *num* bytes
 - Returns actual number of bytes written (between 0 and *num*)

TCP Server

Configuring a Passive-Mode Socket

```
int listen(int sockfd, int backlog);
```

- ❑ Server needs to initialize socket to receive incoming TCP connections
 - ❖ Also called “listening” socket
- ❑ Performed by **listen()**
 - ❖ *sockfd* is the TCP socket (already bound to an address)
 - ❖ *backlog* is the number of incoming connections the kernel should be able to queue for the application
 - ❖ Returns 0 normally, -1 on error

Accepting a Connection

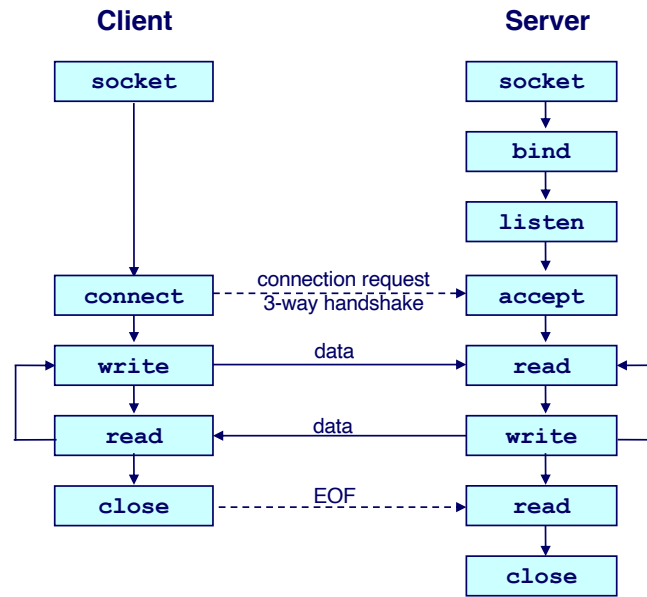
```
int accept(int sockfd,  
           struct sockaddr *cliaddr,  
           socklen_t *addrlen);
```

- ❑ **listen()** must be followed by **accept()** to obtain parameters of accepted connection
 - ❖ *sockfd* is the passive mode TCP socket returned by the `socket()` call
 - ❖ *cliaddr* is a pointer to allocated space for the socket address structure
 - ❖ *addrlen* is a *value-result* argument
 - must be set to the size of *cliaddr* when calling
 - on return, will be set to be the number of used bytes in *cliaddr*

Accepting a Connection (continued)

- ❑ **accept()** returns new a socket descriptor or -1 on error
 - ❖ Different from the original listening socket descriptor
- ❑ Must use the new descriptor to read and write data

Example Sequence of Socket Calls



TCP Client/Server

UDP Client/Server

Sending Data to a Socket

```
ssize_t sendto(int sockfd,  
               const void *buff, size_t nbytes, int flags,  
               const struct sockaddr *to, socklen_t addrlen);
```

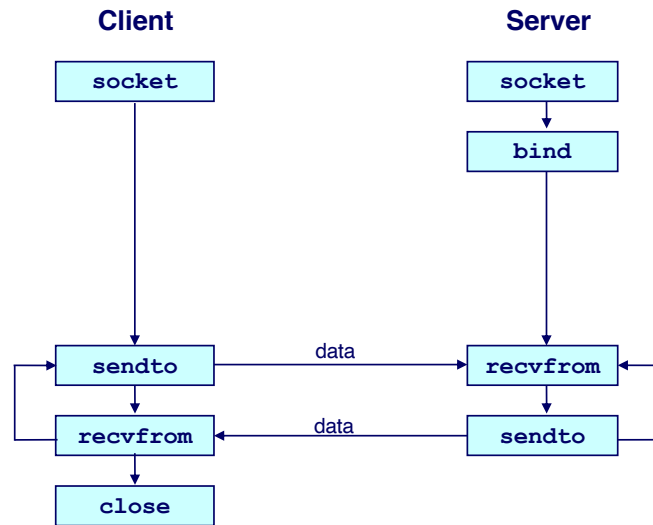
- The first three arguments, *sockfd*, *buff*, and *nbytes*, are identical to the first three arguments for `write()`
- *to* is a socket address structure containing the protocol address of where the data is to be sent
- Writes are non-blocking by default
 - Might not be able to write all *nbytes* bytes
 - Returns actual number of bytes written (between 0 and *nbytes*)

Receiving Data from a Socket

```
ssize_t recvfrom(int sockfd,  
                 void *buff, size_t nbytes, int flags,  
                 struct sockaddr *from, socklen_t *addrlen);
```

- The first three arguments, *sockfd*, *buff*, and *nbytes*, are identical to the first three arguments for `read()`
- The socket address structure of *from* is filled with the address of who sent the datagram, and the size this address structure is also returned in *addrlen*
- By default, **`recvfrom()`** will block until data is available
 - Non-blocking option available
- Returns number of bytes read into the buffer
 - Between 1 and *nbytes* bytes when connection is open
 - 0 indicates EOF (i.e., other side closed connection) with TCP

Example Sequence of Socket Calls



UDP Client/Server