

CSE 156/L

threading pitfalls

Threads

- Main purpose is to enable **concurrency**, thereby improving the performance of the system
 - Reduce latency, i.e., using parallel algorithms
 - Overlap latency, doing other work while waiting for a long-running task to finish
 - Improve throughput by executing multiple copies of the same work/service
- Also, good to take advantage of the available hardware resources, e.g., multiple cores

Thread Correctness

- Race conditions
 - Execution order of access to shared data
 - Non-deterministic program outcome
- Deadlock
 - Take care acquiring **multiple** locks: order locks
 - Use timed non-blocking calls

Thread Safety

- Functions called from a thread must be *thread-safe*
- Four overlapping groups of thread-unsafe functions:
 1. Failing to protect shared variables
 2. Relying on persistent state across invocations
 3. Returning a pointer to a static variable
 4. Calling thread-unsafe functions

Thread-Unsafe Functions

1. Failing to protect shared variables

- Fix: Use lock or semaphore operations
- Issue: Synchronization operations will slow down code, due to overhead and contention

Thread-Unsafe Functions

2. Relying on persistent state across multiple function invocations

- E.g., random number generator relies on static state
- E.g., strtok(), breaking a string into smaller tokens
- Fix: Rewrite function so that caller passes in all necessary state

```
static unsigned int next = 1;

/* rand - return pseudo-random integer on 0..32767 */
int rand(void)
{
    next = next*1103515245 + 12345;
    return (unsigned int)(next/65536) % 32768;
}

/* srand - set seed for rand() */
void srand(unsigned int seed)
{
    next = seed;
}
```

Thread-Unsafe Functions

3. Returning a ptr to a static variable

– Fixes:

a. Rewrite code so caller passes pointer to struct

- Issue: Requires changes in caller and callee

```
struct hostent *  
gethostbyname(char *name)  
{  
    static struct hostent h;  
    //contact DNS and fill in h  
    return &h;  
}
```

```
hostp = malloc(...);  
gethostbyname_r(name, hostp);
```

b. Lock-and-copy

- Issue: Requires only simple changes in caller (and none in callee)
- However, caller must free memory

```
struct hostent *  
gethostbyname_ts(char *p)  
{  
    struct hostent *q = malloc(...);  
    lock(&mutex); /* lock */  
    p = gethostbyname(name);  
    *q = *p;      /* copy */  
    unlock(&mutex);  
    return q;  
}
```

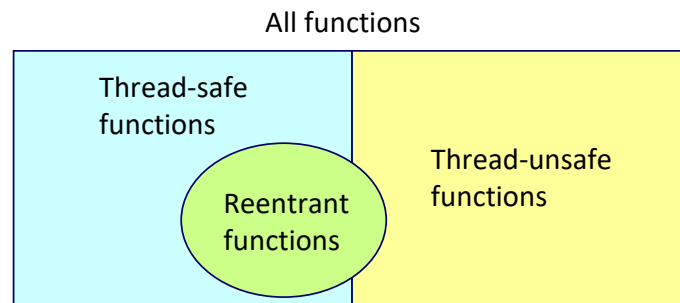
Thread-Unsafe Functions

4. Calling thread-unsafe functions (libraries)

- Calling one thread-unsafe function makes an entire function thread-unsafe
- Fix: Modify the function so it calls only thread-safe functions

Re-entrant Functions

- A function is reentrant if it accesses NO shared variables when called from multiple threads
- Reentrant functions are not a proper subset of the set of thread-safe functions



Thread-Safe Library Functions

- Many functions in the Standard C Library are thread-safe (using synchronization internally)
 - E.g.: `malloc`, `free`, `printf`, `scanf`
- But there are a few exceptions:
 - E.g.: `strtok`, `system`, `strerror`

Thread-unsafe function	Reentrant version
<code>asctime</code>	<code>asctime_r</code>
<code>ctime</code>	<code>ctime_r</code>
<code>gethostbyaddr</code>	<code>gethostbyaddr_r</code>
<code>gethostbyname</code>	<code>gethostbyname_r</code>
<code>inet_ntoa</code>	(none)
<code>localtime</code>	<code>localtime_r</code>
<code>rand</code>	<code>rand_r</code>

Thread Bugs

- Deadlock (e.g., dining philosopher's problem)
 - If a thread needs more than one mutex, follow a strict ordering
 - May use mutex back-off loop to acquire all locks
- Priority inversion
 - Need three threads with different priorities
 - Lower priority thread prevents high priority thread from running

Threads Summary

- Provide another mechanism for writing concurrent programs
 - Cheaper overhead than processes
- Easy to share data between threads
- Risk of race conditions and deadlock
- Cost of easy sharing:
 - Error-prone synchronization code within a process
 - E.g., access to any global variable