


# I/O Multiplexing

## What is I/O Multiplexing?

- When an application needs to handle multiple I/O descriptors at the same time
  - E.g., file and socket descriptors, or multiple socket descriptors
- Application processes both *interactive* input and *network* socket
- Server handles multiple ports and protocols
  - Server handles both TCP listen and connected sockets
  - Server handles both TCP and UDP
- When I/O on any one descriptor can result in *blocking*

# Non-Forking Concurrent Server



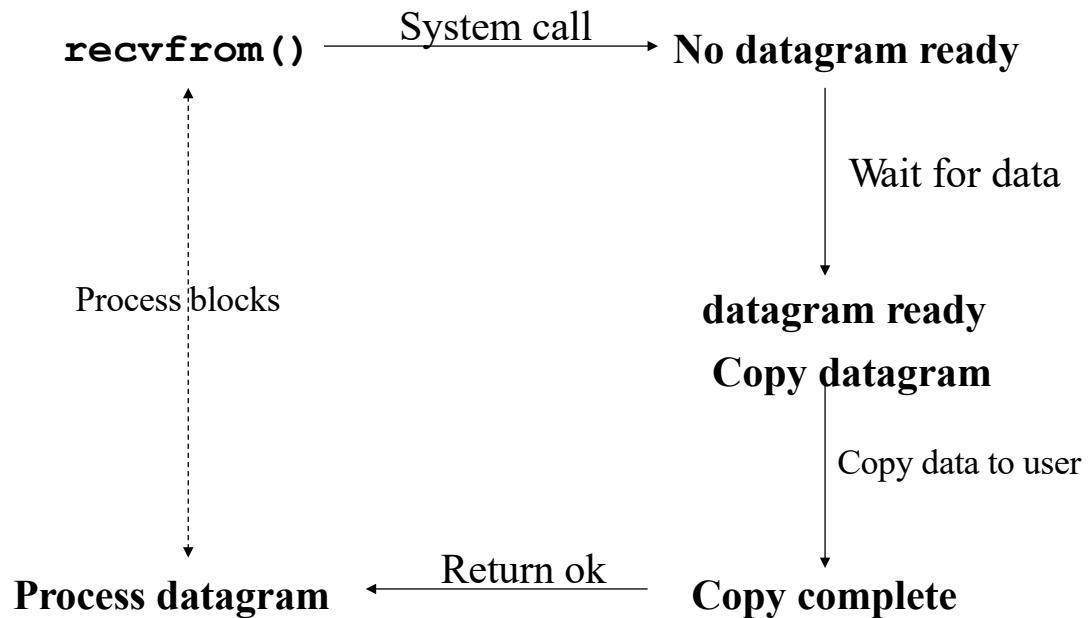
## I/O models

1. Blocking I/O
2. Non-blocking I/O
3. I/O multiplexing – `select()`, `poll()`
4. Signal driven I/O
5. Asynchronous I/O

# Blocking I/O

Application

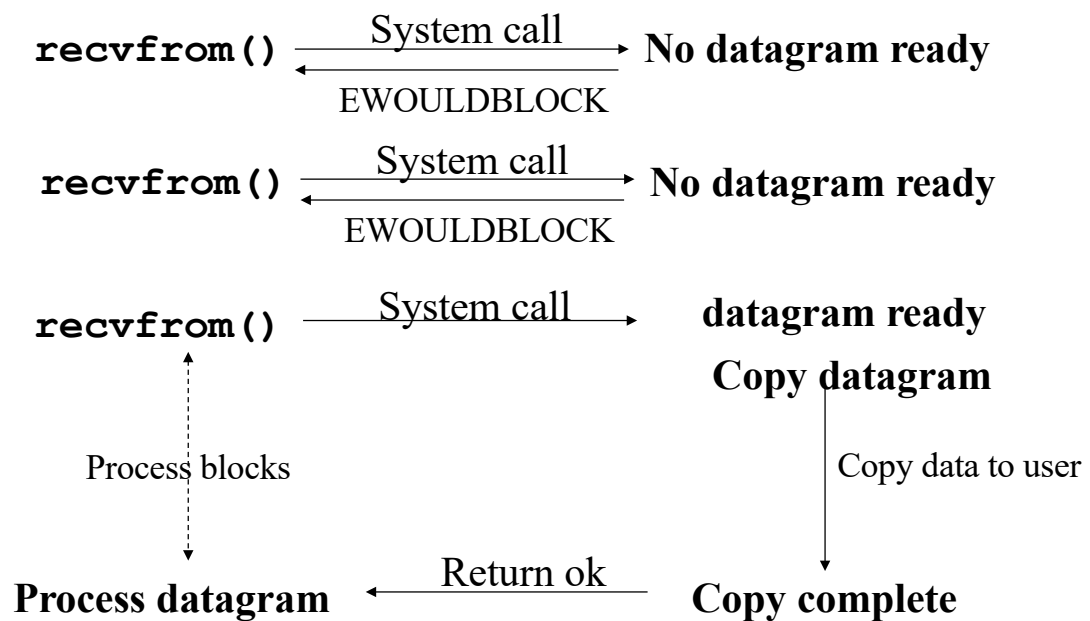
Operating system



# Non-Blocking I/O

Application

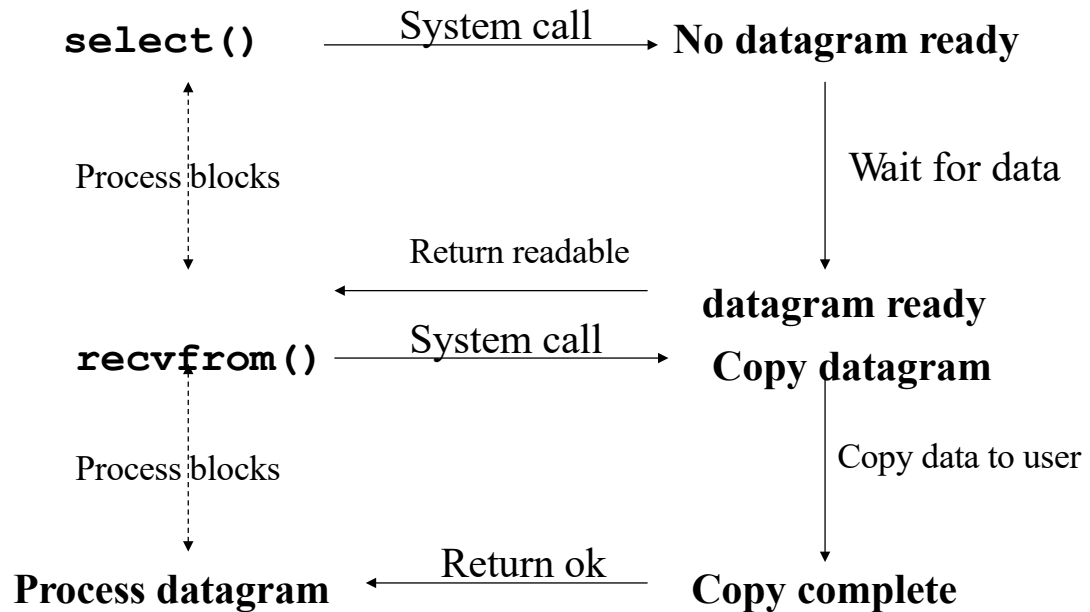
Operating system



# I/O Multiplexing

Application

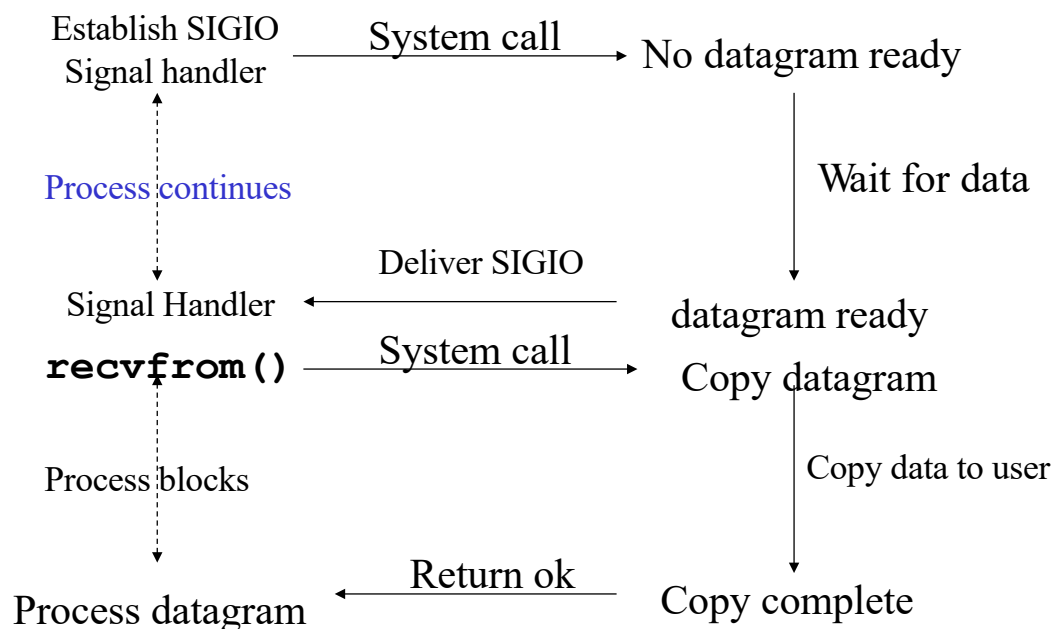
Operating system



# Signal driven I/O

Application

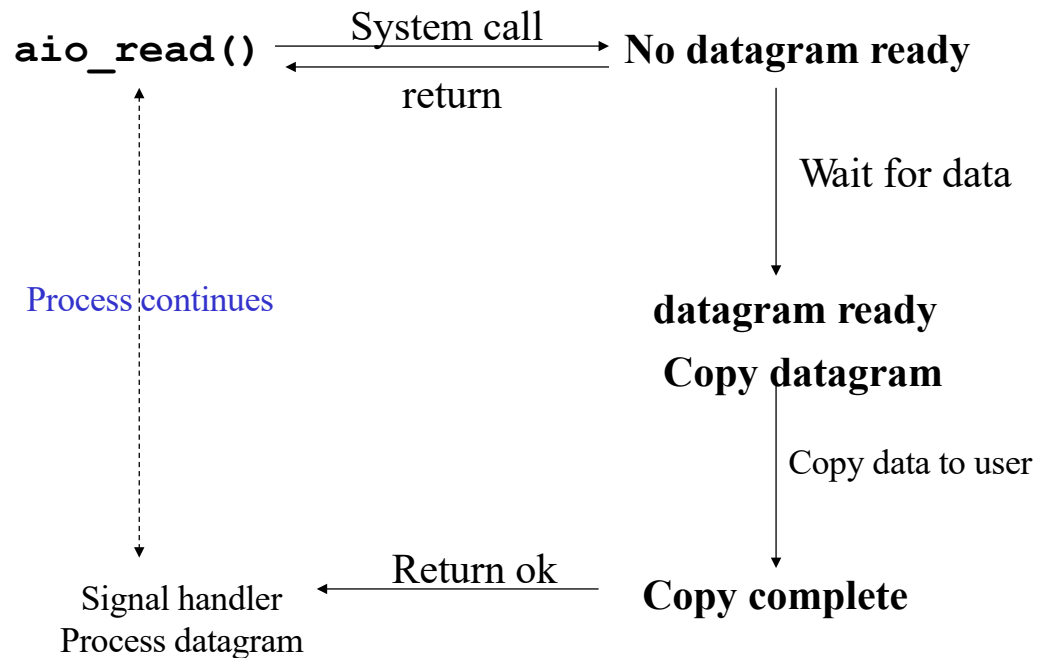
Operating system



# Asynchronous I/O

Application

Operating system



## select() call

- Allows a process to *wait* for an event to occur on any one of its descriptors:
  - Wait forever
  - Wait for a fixed amount of time
  - Do not wait at all! This is called *polling*
- Types of *events*
  - Ready for read
  - Ready for write
  - Exception condition

# select() call

```
int  select(int maxfdp1,          /* max. fd + 1 */
            fd_set *readfds,      /* read ready? */
            fd_set *writefds,     /* write ready? */
            fd_set *exceptfds,    /* exception? */
            struct timeval *timeout);
```

Returns positive count of ready descriptors, 0 on timeout, -1 on error

```
struct timeval {
    long tv_sec;  /* seconds */
    long tv_usec; /* microseconds */
};
```

## fd\_set

- Describes a set of descriptors that we want to wait on for events
- Traditionally held 1024 descriptors, but should be configurable for newer kernels
- Defines manipulation macros

```
void FD_ZERO(fd_set *fds)
void FD_SET(int fd, fd_set *fds)
void FD_CLR(int fd, fd_set *fds)
int  FD_ISSET(int fd, fd_set *fds)
```

## Value-result arguments

- Select modifies descriptor sets pointed to by *readset*, *writeset*, and *exceptset* pointers
- On function call
  - Specify value of descriptors that we are interested in
- On function return
  - Result indicates which descriptors are ready
- Use FD\_ISSET macro on return to test a specific descriptor in an fd\_set structure
  - Any descriptor not ready will have its bit cleared
  - You need to turn on all the bits in which you are interested on all the descriptor sets each time you call select

## maxfdp1 argument

- Specifies the maximum number of descriptors to be tested
- Its value is the maximum descriptor plus one
  - Descriptors 0, 1, 2, up through and including *maxfdp1*-1 are tested
  - Calculate the maxfdp1 value based on the current descriptor set *for efficiency* of processing fd\_set
- Constant FD\_SETSIZE defined by including header file <sys/select.h>
  - It is the number of descriptors in the fd\_set datatype (traditionally 1024)

## Socket conditions

- Handling conditions where peer can send data, RST or FIN




Fig. 6.8

## Socket conditions (2)

- Can handle conditions where peer can send data, RST or FIN
- Peer TCP sends data
  - The socket becomes readable and `read` returns greater than 0 (number of bytes of data)
- Peer TCP sends a FIN (terminate or close)
  - The socket becomes readable and `read` returns 0 (i.e., EOF)
- Peer TCP sends a RST (crashed and rebooted)
  - The socket becomes readable and returns -1
  - **errno** contains the specific error code



# Ready conditions

Summary of conditions that cause a socket to be ready for select()

Condition	Readable?	Writable?	Exception?
Data to read	•		
Read-half of the connection closed	•		
New connection ready for listening socket	•		
Space available for writing		•	
Write-half of the connection closed		•	
Pending error	•	•	
TCP out-of-band data			•

## Non-forking concurrent server

```
fdset rdset, wrset;
int listenfd, connfd1, connfd2;
int maxfdp1;

/* do connection establishment etc */
.....

/* initialize */
FD_ZERO(&rdset);
FD_ZERO(&wrset);
```

```

for ( ;; ) {
    FD_SET(connfd1, &rdset);
    FD_SET(connfd2, &wrset);
    FD_SET(listenfd, &rdset);

    maxfdp1 = max(connfd1, connfd2, listenfd) + 1;
    /* wait for some event */
    Select(maxfdp1, &rdset, &wrset, NULL, NULL);

    if (FD_ISSET(connfd1, &rdset)) {
        /* read data from connfd */
    }
    if (FD_ISSET(connfd2, &wrset)) {
        /* write data to connfd */
    }
    if (FD_ISSET(listenfd, &rdset)) {
        /* process a new connection */
    }
}

```

## High resolution timer!

- Select can be used as a millisecond resolution timer.

```
select(0, NULL, NULL, NULL, &timeval);
```

- Usual `sleep()` call has resolution of seconds.

## Fig 6.13 (strcliselect02.c)

```

1 #include "unp.h"
2 void
3 str_cli(FILE *fp, int sockfd)
4 {
5     int maxfdp1, stndeof;
6     fd_set rset;
7     char buf[MAXLINE];
8     int n;
9     stndeof = 0;
10    FD_ZERO(&rset);
11    for (;;) {
12        if (stndeof == 0)
13            FD_SET(fileno(fp), &rset);
14        FD_SET(sockfd, &rset);
15        maxfdp1 = max(fileno(fp), sockfd) + 1;
16        Select(maxfdp1, &rset, NULL, NULL, NULL);
17
18        if (FD_ISSET(sockfd, &rset)) { /* socket is readable */
19            if ((n = Read(sockfd, buf, MAXLINE)) == 0) {
20                if (stndeof == 1)
21                    return; /* normal termination */
22                else
23                    err_quit("str_cli: server terminated prematurely");
24            }
25            Write(fileno(stdout), buf, n);
26        }
27        if (FD_ISSET(fileno(fp), &rset)) { /* input is readable */
28            if ((n = Read(fileno(fp), buf, MAXLINE)) == 0) {
29                stndeof = 1;
30                Shutdown(sockfd, SHUT_WR); /* send FIN */
31                FD_CLR(fileno(fp), &rset);
32                continue;
33            }
34            Writen(sockfd, buf, n);
35        }
36    }
37 }

```

## Multiple client descriptors (e.g.)

- Instead of fork() per client, use select()
- After the second client connection is established (assuming connected descriptor returned by accept is 5)

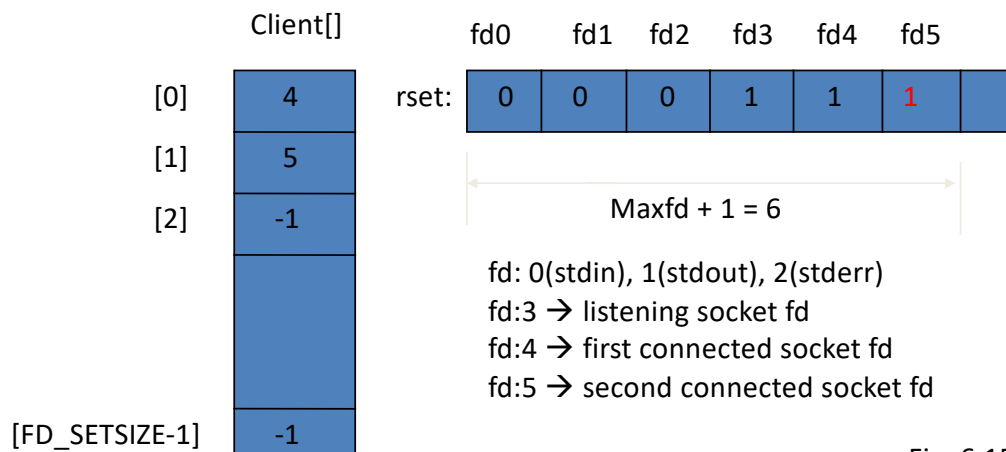


Fig. 6.15

## Fig. 6.22 (tcpselect01.c)

```
25 for (;;) {
26     rset = allset; /* structure assignment */
27     nready = Select(maxfd + 1, &rset, NULL, NULL, NULL);

28     if (FD_ISSET(listenfd, &rset)) { /* new client connection */
29         clien = sizeof(cliaddr);
30         connfd = Accept(listenfd, (SA *) &cliaddr, &clilen);

31         for (i = 0; i < FD_SETSIZE; i++)
32             if (client[i] < 0) {
33                 client[i] = connfd; /* save descriptor */
34                 break;
35             }
36         if (i == FD_SETSIZE)
37             err_quit("too many clients");
38         FD_SET(connfd, &allset); /* add new descriptor to set */
39         if (connfd > maxfd)
40             maxfd = connfd; /* for select */
41         if (i > maxi)
42             maxi = i; /* max index in client[] array */

43         if (--nready <= 0)
44             continue; /* no more readable descriptors */
45     }
46     for (i = 0; i <= maxi; i++) { /* check all clients for data */
47         if (sockfd = client[i] < 0)
48             continue;
49         if (FD_ISSET(sockfd, &rset)) {
50             if ((n = Read(sockfd, buf, MAXLINE)) == 0) {
51                 /* connection closed by client */
52                 Close(sockfd);
53                 FD_CLR(sockfd, &allset);
54                 client[i] = -1;
55             } else
56                 Writen(sockfd, buf, n);

57         if (--nready <= 0)
58             break; /* no more readable descriptors */
59     }
60 }
61 }
62 }
```

## pselect()

- Const nanosec timespec for timeout parameter
- Set a blocking mask while waiting, restore original mask when returning

```
int pselect(int maxfdp1,
            fd_set *readset,
            fd_set *writeset,
            fd_set *exceptset,
            const struct timespec *timeout,
            const sigset_t *sigmask);
```

```
struct timespec {
    time_t tv_sec; /* seconds (long) */
    long tv_nsec; /* nanoseconds */
};
```

## pselect()

Race condition can exist between the call to select() and signal delivery

```
if (intr_flag)
    handle_intr();    /* handle the signal */

if ( (nready = select( ... )) < 0 ) {
    if (errno == EINTR) {
        if (intr_flag)
            handle_intr();
    }
    ...
}
```

← ??

## pselect()

Signal mask argument overrides the existing signal mask

```
sigset_t newmask, oldmask, zeromask;

sigemptyset(&zeromask);
sigemptyset(&newmask);
sigaddset(&newmask, SIGINT);

sigprocmask(SIG_BLOCK, &newmask, &oldmask);    /* block SIGINT */
if (intr_flag)
    handle_intr();    /* handle the signal */
if ( (nready = pselect ( ... , &zeromask)) < 0 ) {
    if (errno == EINTR) {
        if (intr_flag)
            handle_intr ();
    }
    ...
}
```

# poll()

- Provides similar functionality to select()
- Events encoded in bits specifying certain conditions
- Wait forever INFTIM (-1 or < 0); or don't block 0; or wait for > 0

```
#include <poll.h>
```

```
int poll(struct pollfd *fdarray,  
         unsigned long nfd, /* array size */  
         int timeout);      /* msec to wait */
```

```
struct pollfd {  
    int fd; /* descriptor to check */  
    short events; /* events of interest */  
    short revents; /* events occurred */  
};
```

## Poll() Parameters

- Have two variables per descriptor
  - Input value events member **event**
  - Returns status for a descriptor in the corresponding **revents** member

	Constant	Input to <i>events</i> ?	Result from <i>revents</i> ?	Description
inputs	POLLIN	•	•	Normal or priority band data can be read
	POLLRDNORM	•	•	Normal data can be read
	POLLRDBAND	•	•	Priority band data can be read
	POLLPRI	•	•	High-priority data can be read
outputs	POLLOUT	•	•	Normal data can be written
	POLLWRNORM	•	•	Normal data can be written
	POLLWRBAND	•	•	Priority band data can be written
errors	POLLERR		•	Error has occurred
	POLLHUP		•	Hangup has occurred
	POLLNVAL		•	Descriptor is not an open file

- Set **fd** member < 0 to ignore it
- Note fdarray allocated by caller; no fixed-size data type similar to fd\_set

## Poll Conditions (e.g.)

- Three classes of data identified:
  - Normal, priority, high-priority
- Some events for backward compatibility, e.g., POLLIN and POLLOUT
  - POLLIN: data other than high-priority data may be read without blocking
  - POLLRDNORM: normal data (priority band equals 0) may be read without blocking
  - POLLWRNORM: same as POLLOUT
- Regular TCP/UDP considered normal
- TCP out-of-band is priority band
- TCP read half-close is normal
- TCP error can be either normal or error (POLLERR)

\*\*\*